

Programming Assignment 1: Programming with Sockets

In this assignment you will be asked to implement an HTTP client and server running a pared down version of HTTP/1.0. You will have to extend the client and server to make use of some of the application level headers we have studied in class. For those of you who already have socket programming experience, the basic part of this assignment will be fairly easy. I am therefore adding an extra credit component for those students who want a more challenging assignment. This project can be completed in either C/C++, Java, or other languages you prefer.

What to Turn In

When you hand in your programming assignment, you should include:

- A program listing containing in-line documentation. Uncommented code will be heavily penalized.
- A separate (typed) document of a page or so describing the overall program design, a verbal description of "how it works", and design tradeoffs considered and made. Also describe possible improvements and extensions to your program (and sketch how they might be made).
- A separate description of the test cases you ran on your program to convince yourself (and me) that it is indeed correct, and execution traces showing these test being run. Also describe any cases for which your program is known not to work correctly. The test cases you should hand are described below.

The HTTP Client

Your client should take command line arguments specifying a server name or IP address, the port on which to contact the server, the method you use, and the path of the requested object on the server. You are going to implement two methods of HTTP: `GET` and `PUT`.

- `GET`
The format of the command line is

```
myclient host port_number GET filename
```

The basic client action should proceed as follows:

1. Connect to the server via a connection-oriented socket.
2. Submit a valid HTTP/1.0 GET request for the supplied URL.

```
GET /index.html HTTP/1.0
```

(end with extra CR/LF)

3. Read (from the socket) the server's response and display it as program output.

Once you have this part of the client working, you should demonstrate it with the following two test cases:

1. Use it to get a file of your choosing from a "real" web server on the internet. For example,

```
myclient www.cnn.com 80 GET index.html
```

2. Use it to get a file from your own server program. For example, your server is running on pc1.cs.uml.edu, port number 5678.

```
myclient pc1.cs.uml.edu 5678 GET index.html
```

This command would result in an HTTP GET request to pc1.cs.uml.edu for index.html on port 5678, and get the file index.html back to the client.

- PUT

The format of the command line is

```
myclient host port_number PUT filename
```

The basic client action should proceed as follows:

1. Connect to the server via a connection-oriented socket.
2. Submit a PUT request for the supplied file:

PUT filename
extra CR/LF.

(Once your server program receives such a request, it should expect to receive the file and save it to disk.)

3. Send the file to the server.
4. Wait for server's reply.

Once you have this part of the client working, you should test it with your own server: send out a file to your server, the server should save the file and send back a response.

The HTTP Server

Your server should take command line arguments specifying a port number. For example,

```
myserver 5678
```

The basic server action should proceed as follows

1. Initialize the server.
2. Wait for a client connection on the port number specified by command line argument.
3. When a client connection is accepted, read the HTTP request.
4. Construct a valid HTTP response including status line, any headers you feel are appropriate, and, of course, the requested file in the response body.

For GET , if the server receives the "GET index.html HTTP/1.0" request, it sends out "200 OK" to the client, followed by the file index.html. If the requested file doesn't exist, the server sends out "404 Not Found" response to the client.

For PUT , if the server receives the "PUT test.txt" request, it will save the file as test.txt. If the received file from client is successfully created, the server sends back a "200 OK File Created" response to the client.

5. Close the client connection and loop back to wait for the next client connection to arrive.

Notice that your server will be structured around an infinite loop. That means that you must interrupt your server with a termination signal to stop it. Make sure your server code shuts down gracefully when terminated. That means closing any open sockets, freeing allocated memory, etc.

Once you get your server working, demonstrate it with the following two test cases:

First, use an ordinary browser such as Netscape or Internet Explorer to get a html file from your server. For example, your server is running at host `pc1.cs.uml.edu` on port number 5678, and there is a file `index.html` in the current directory. In the URL box of the web browser, type in `pc1.cs.uml.edu:5678/index.html`, the browser should fetch the file and display it.

Second, use your own client to get a file.

Third, use your own client to put a file.

Programming notes

Here are a few tips/thoughts to help you with the assignment:

- You must chose a server port number greater than 1023 (to be safe, choose a server port number larger than 5000).
- I would strongly suggest that everyone begin by writing a client and getting its test cases to work. Then write a sequential server and test it with your client. Finally, write a concurrent server, if you want to do the extra credit portion of the assignment.
- In writing your code, make sure to check for an error return from your system calls or method invocations, and display an appropriate message. In Java, this means using `IOException()`; in C this means checking and handling error return codes from your system calls. See the documentation noted above.
- Most of you will be running both the client and server on the same machine (e.g., by starting up the server and running it in the background, and then starting the client) under Unix. Recall the use of the ampersand to start a process in the background. If you need to kill your server after you have started it, you can use the UNIX *kill* command. Use the UNIX *ps* command to find the process id of your server
- Make sure you close every socket that you use in your program. If you abort your program, the socket may still hang around and the next time you try and bind a new socket to the port ID you previously used (but never closed), you may get an error. Also, please be aware that port ID's, when bound to sockets, are system-wide values and thus other students may be using the port number you are trying to use. On UNIX systems, you can run the command "netstat" to see which port numbers are currently assigned.

Extra Credit: Concurrent Server [+ 10%]

Rewrite your server to be a concurrent server - that is a server that waits on the welcoming socket and then creates a new thread or process to handle the incoming request. If you're doing the assignment in Java, see the section "Supporting Multiple Clients" in Sun's socket programming tutorial,

<http://java.sun.com/docs/books/tutorial/networking/sockets/clientServer.html>.